

- JupyterLab Desktop
- JupyterLab
- VS Code, Positron, etc.

Please type and execute this code:

```
import pandas as pd  
print(pd.__version__)
```

'2.3.3'

← you should get
something like this



jupyter

The Jupyter logo consists of a large orange circle with three smaller gray circles positioned above, below, and to the right of it.

Raise your hand if you cannot import
Pandas or if you get a version < 2.0

Roll Call

- Your name
- Role or class year
- Department or major
- A source or form of data you regularly use or want to learn to use

ICPSR

United States®
Census Bureau

eurostat 



Meet Your Instructor



- Erin M. Winter, she/her
- Data Management Librarian
- MS Library Science 2023, MS Computer Science 2017
- Supports Python, HPC, GitHub, and data management at Data Services

uolib.link/data

Workshop Policies: Attendance



Come to Knight 267B
to attend in person.

OR

zoom

Link automatically
sent to the *email you
registered with* 3
hours before the
workshop

Lecture Notes and Slides on GitHub

The screenshot shows a GitHub repository interface. At the top, it says 'pandas-workshop' (Public). Below that, there are buttons for 'Unpin' and 'Watch'. The main navigation bar includes 'main' (selected), '1 Branch', '0 Tags', a search bar ('Go to file'), and buttons for 'Add file' and 'Code'. A commit history for the 'final' branch is displayed, starting with a commit from 'winter' (9 months ago) that updated the repository's structure. Other commits include 'final' updates for 'Notebooks' and 'Slides', and resource-related changes like 'day one resources' and 'resources'. All commits are dated 9 months ago.

File / Commit Message	Author	Date
Archive/F24	winter	update structure
Notebooks	winter	final
Slides	winter	final
.gitignore	winter	day one resources
README.md	winter	update structure
environment.yml	winter	update
resources.md	winter	resources

GitHub Page

Notebooks

Slides

Updated a few days
after the workshop

uolib.link/pandas-notes

Prerequisites

- Software Carpentries:
Plotting and Programming in
Python

OR

- JupyterLab interface
- **variables, assignment**
- standard Python primitives
 - int, float, string, boolean
- control flow: if and else
- *functions, scope*



You will get opportunities to practice these if you're still grasping them.

Workshop Objectives

Goals

- Manipulate, analyze, and visualize data in Pandas and related libraries
- Import data from a variety of sources
- Identify and correct common problems in tabular data
- Learn basic data visualization

Not Covered

- Topics in “big data”
- Math
- Non-tabular data: ie. genetic sequences, geodatabases, images

Why Pandas?

- Free, easy to learn
- Compatible with a growing data science ecosystem in Python
 - Visualization
 - Machine learning
 - Statistics
 - Scientific applications
- Flexible in terms of input and output



bokeh



pandas



Data Wrangling

From Raw Data to Analysis

- Handle missing values
- Convert measurements, apply functions, and correct mistakes at scale
- Prepare data to be used for other applications
- Create reusable code so you don't have to do it from scratch again

I need to use this data but this variable is in the wrong format.


A professor gave me this data but it won't open.

Someone else entered the data incorrectly.

I want to put data from multiple sources together.


Creating DataFrame Objects

- **Can be constructed directly from Python objects like dictionaries**
- **Loaded from a file or URL:**
 - delimited text file (.txt, .tsv)
 - comma-separated values (.csv)
 - Excel sheet (.xlsx)
 - JSON dictionary (.json)
 - Stata/SAS files (.sav, .sas, .dta)



Getting Started

- Create a folder for your project named **pandas-workshop**
- Open the folder in JupyterLab
- Create a subfolder for your data called ***data***



You will be
downloading your data
to this folder.

A Little Modern Art


- Our first dataset is the MOMA's Watson Library Index of Asian American and Pacific Islander Artists.
- This is a collection of AAPI artists featured in the Watson Library's catalog of exhibit records and artist biographies.



Kamekichi Tokita
(1897-1948)



Tokita, Kamekichi. Alley. 1929. Seattle Art Museum. Collection of Shokichi & Elise Y. Tokita.



Loading Data From A File

- Find the download link in the *resources.md* in the class GitHub page
- Download this file from the Watson Library GitHub by clicking the download button
- Copy this *.csv* file into your JupyterLab project directorys

Reading Input

Reading files into Pandas is as simple as **matching the file type to the name of a function...**

- pd.read_csv()
- pd.read_json()
- pd.read_excel()
- pd.read_sql()
- pd.read_stata()
- ...

Format	Type	Data Description	Reader	Writer
	text	CSV	<code>read_csv</code>	to_csv
	text	Fixed-Width Text File	<code>read_fwf</code>	NA
	text	JSON	<code>read_json</code>	to_json
	text	HTML	<code>read_html</code>	to_html
	text	LaTeX	<code>Styler.to_latex</code>	NA
	text	XML	<code>read_xml</code>	to_xml
	text	Local clipboard	<code>read_clipboard</code>	to_clipboard
	binary	MS Excel	<code>read_excel</code>	to_excel
	binary	OpenDocument	<code>read_excel</code>	NA
	binary	HDF5 Format	<code>read_hdf</code>	to_hdf


https://pandas.pydata.org/docs/dev/user_guide/io.html

Understanding Your Data (I/O)

After reading from a file to a DataFrame, check

- delimiting (are the columns separated?)
- the shape – a (row #, column #) tuple
- the types of each variable – dtypes
- which columns have nulls/nones/NaNs
(and **why**)

?



If it is **not** what you expect, you probably need to adjust the formatting arguments passed to Pandas.

What Can Go in a DataFrame?

name	colour	location	seed	shape	sweetness	water_content
apple	red	canada	TRUE	round	TRUE	84

int64, float64

bool

object, Categorical

datetime64

Integers, Floats

Booleans

Strings

Dates

Columns with values of a type that Pandas cannot infer will be labeled as objects.

The Pandas DataFrame

Each **DataFrame** column is a **Series**.

.index


0
1
2
3
4

	name	colour	location	seed	shape	sweetness	water_content	weight
0	apple	red	canada	TRUE	round	TRUE	84	100
1	banana	yellow	mexico	FALSE	long	TRUE	75	120
2	cantaloupe	orange	spain	TRUE	round	TRUE	90	1360
3	dragon fruit	magenta	china	TRUE	round	FALSE	96	600
4	elderberry	purple	austria	FALSE	round	TRUE	80	5

.columns

axis = 1 →

axis = 0



All values in the DataFrame have an indexed position in `.iloc[row_index, column_number]` form.

Total elements = `.size`
R, C form = `.shape`



The Pandas Series

The **Series** is the simplest of the Pandas data structures: a one-dimensional array with an index. The alpha-numeric index is exactly as long as the data.

.index

0	1	2	3	4	5	6	7
burgundy	red	green	gray	blue	yellow	orange	teal

All values in the Series have an indexed position in
[index] or .iloc[index_num] form.

.index

a	b	c	d	e	f	g	h
burgundy	red	green	gray	blue	yellow	orange	teal

Slicing and Indexing with .loc

		Column Labels		
		name	colour	location
Index	a	apple	red	canada
	b	banana	yellow	mexico
	c	cantaloupe	orange	spain
	d	dragon fruit	magenta	china
	e	elderberry	purple	austria

.loc slicing is inclusive on both ends because it is intended to behave like R (not Python)

.loc – index and column based slicing

.loc['a'] returns a **pandas.Series** with the first row

.loc['a', 'location'] returns the string “canada”

.loc['b':'d', 'name': 'colour'] returns the DataFrame below

	name	colour
b	banana	yellow
c	cantaloupe	orange
d	dragon fruit	magenta

Slicing and Indexing with .iloc

		Column Labels		
		name	colour	location
Index	a	apple	red	canada
	b	banana	yellow	mexico
	c	cantaloupe	orange	spain
	d	dragon fruit	magenta	china
	e	elderberry	purple	austria

.iloc – integer-based slicing

.iloc[0] returns a **pandas.Series** with the first row

.iloc[1, 2] returns the string “mexico”

.iloc[2:, 1:3] returns the DataFrame below

.iloc slices behave like Python list slices (exclusive at the end of a range)

	colour	location
c	orange	spain
c	magenta	china
d	purple	austria

Boolean Indexing

- Pandas .loc and .iloc can return subsets of a DataFrame specified by a boolean array.
 - Boolean arrays filter DataFrames or Series by testing values against a condition. NaNs evaluate to *False*.

```
under21 = students['age'] < 21
```

a series of a pd.DataFrame numeric column the “test”, must size (# of rows in students) evaluate to True or False

```
studentsUnder21 = students[under21]
```

only the rows in the students DataFrame where the column ‘age’ < 21

x['col_name'] is shorthand for x.loc[:, 'col_name']

Boolean Indexing with Multiple Conditions

- We can combine boolean indexing with multiple conditions as follows:

```
a pd.DataFrame column test boolean  
ok_to_drink = (students['age'] >= 21) & (operator  
    students['has_id'] == 'Yes')  
test
```

a series of
size (# of
rows in
students)

```
studentDrinkers = students[ok_to_drink]
```

a DataFrame with only the rows in the
students DataFrame where the column 'age'
> 21 AND column 'has_id' has a 'Yes' value

What is a NaN? (null, None, etc.)

The np.NaN is the way Pandas represents **missing values** by default, but missing values occur in almost **all** domains:


- Product of the data entry process
- Missing from the source data
- Measurement error
- Participant declined to respond
- Something went wrong computationally when creating the data

Why Dtypes Matter

continuous	discrete	discrete	ordinal/continuous
Integers, Floats	Booleans	Strings	Dates
sort/compare	filter	count	sort/compare
arithmetic	count	compare	add/subtract
percentiles	sort/compare	alphabetize	filter
filter		filter	count
count			


Inferring New Columns

- Infer a new column from an existing column or columns
 - Map()
 - Vectorized operations
 - $F[\text{mpg}] = F[\text{miles_per_tank}]/F[\text{gallons_per_tank}]$



The Histogram

Every bin of frequencies has an associated count




X, Y - one continuous variable
Y = frequency of the values within each bin

Gives you a sense of the shape of the distribution, very sensitive to the number of bins.


The Bar Chart

Every bar has a category X and a value Y




The Line Graph

Every line has...
X, Y coordinates




The Scatter Plot




Combining Data Sources


- The simplest (and most naïve approach) to combining DataFrames is to **concatenate them**



+




=




Pandas.concat


Axis = index



+



=




Axis = columns

JSON

A structured text-formatting standard compatible with Python, R, JavaScript etc.

Python has a robust built-in JSON parser


JSONs can be treated as dictionaries or lists of dictionaries



From JSON to Pandas DataFrame

A JSON file can be treated as a **list of dictionaries** with a shared schema

For simple JSON files, Pandas can directly read JSON files from URLs or from disk into DataFrames



Merge (Join)



Allows you to join two or more DataFrames using shared values in a shared **key** column to align the rows in the DataFrames

Akin to a SQL join, R's merge, or an Excel VLOOKUP

**Thank you for
attending!**