

- JupyterLab
- JupyterNotebook
  - or an IDE of your choice...

Please type and execute this code:

```
import pandas as pd  
print(pd.__version__)
```

'2.2.3'

← you should get something like this



Raise your hand if you cannot import Pandas **or if you get a version < 2.0**

# Roll Call

- Your name
- Role or class year
- Department or major
- A source or form of tabular data you regularly use or want to learn to use

ICPSR

United States<sup>®</sup>  
**Census**  
Bureau

eurostat 



# Workshop Policies

Lecture Notes  
(.ipynb)

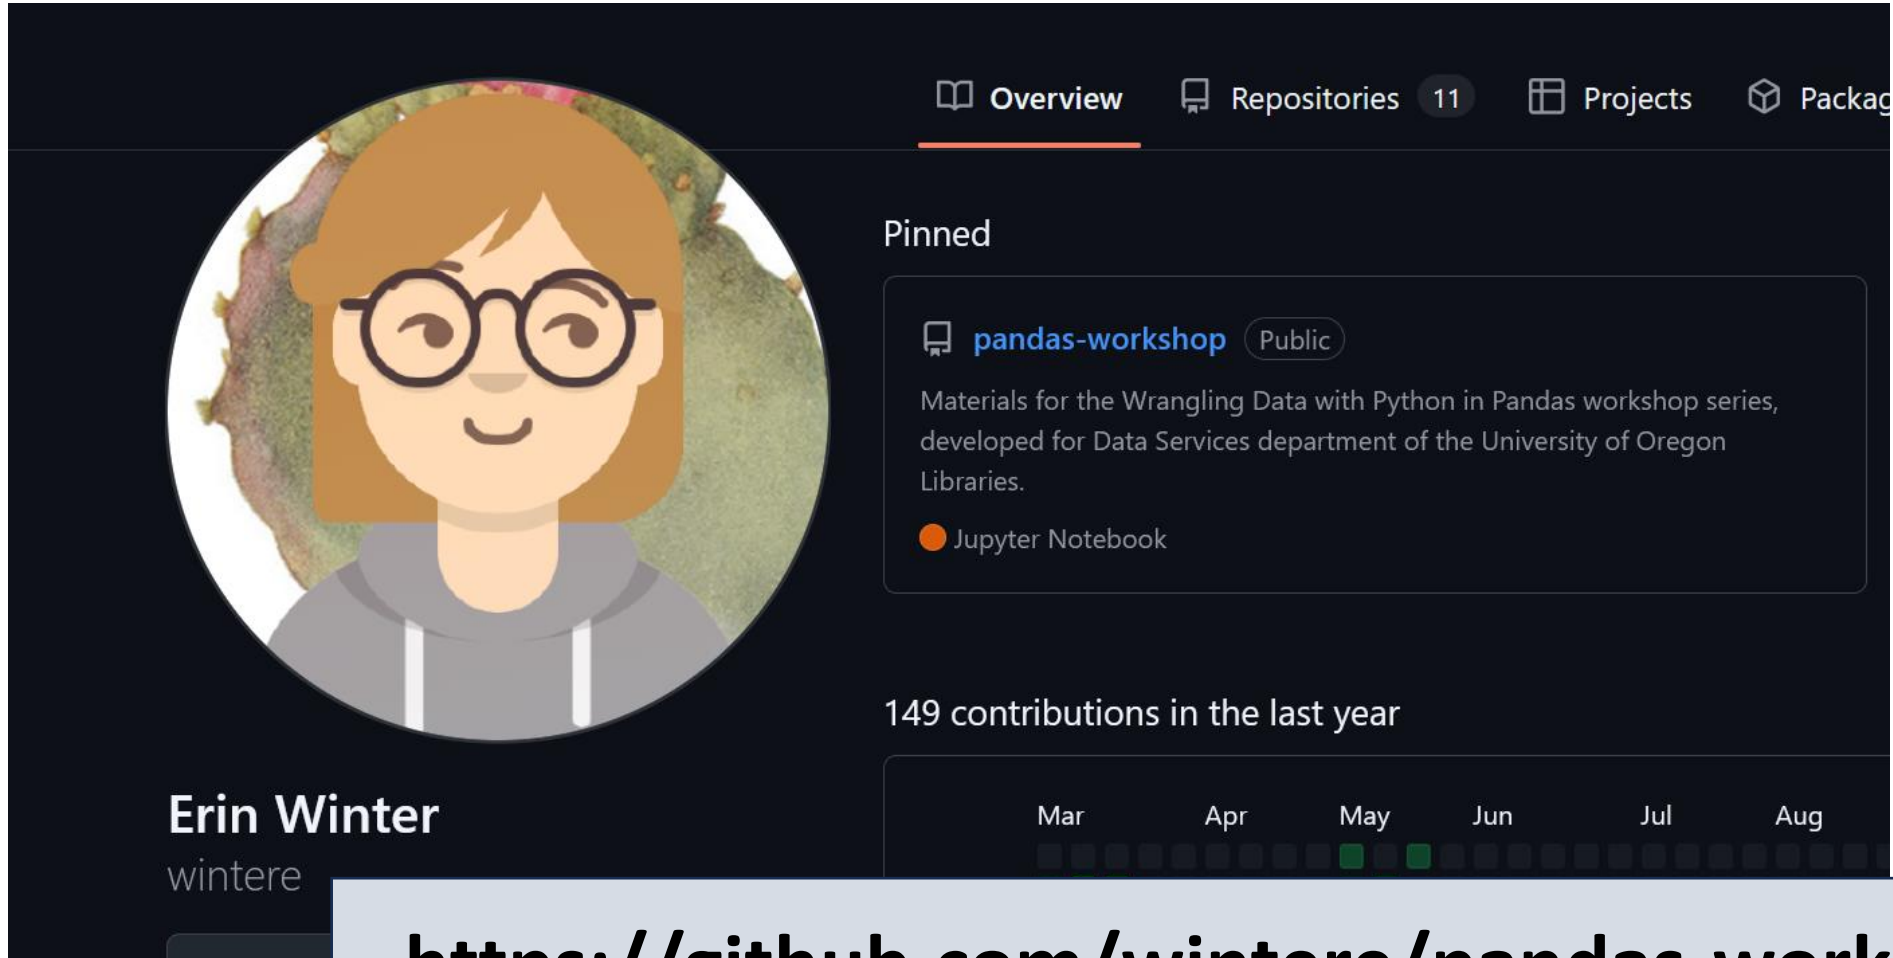
Slides

Updated a few days  
after the workshop



Link automatically  
sent to the *email you*  
*registered with* 24  
hours before the  
workshop


# Lecture Notes and Slides on GitHub




The screenshot shows the GitHub profile of Erin Winter. On the left is a circular profile picture of a woman with brown hair and glasses. Below it, the name "Erin Winter" and the username "wintere" are visible. To the right, the navigation bar includes "Overview" (selected), "Repositories" (11), "Projects", and "Packages". The "Pinned" section features a repository named "pandas-workshop" with a "Public" label. The description states: "Materials for the Wrangling Data with Python in Pandas workshop series, developed for Data Services department of the University of Oregon Libraries." Below this, it says "Jupyter Notebook" with an orange dot icon. At the bottom, it shows "149 contributions in the last year" and a calendar view for the months of March, April, May, June, July, and August, with May and June highlighted in green.

Overview Repositories 11 Projects Packages

Pinned

 **pandas-workshop** Public

Materials for the Wrangling Data with Python in Pandas workshop series, developed for Data Services department of the University of Oregon Libraries.

 Jupyter Notebook

149 contributions in the last year

Mar Apr May Jun Jul Aug

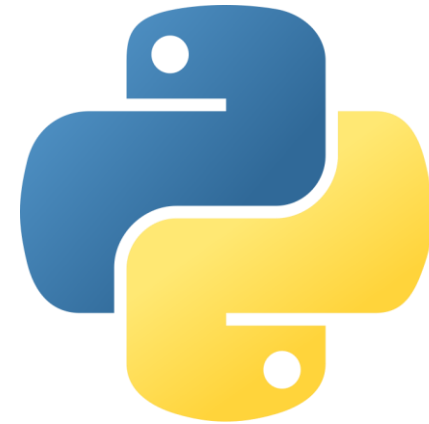
<https://github.com/wintere/pandas-workshop>

# Prerequisites

- Software Carpentries:  
Plotting and Programming in  
Python

OR

- **variables, assignment**
- standard Python primitives
  - int, float, string, boolean
- control flow: if and else
- *functions, scope*



You will get opportunities to practice these if you're still grasping them.

# Workshop Objectives

## Goals

- Manipulate, analyze, and visualize data in Pandas and related libraries
- Import data from a variety of sources
- Identify and correct common problems in tabular data
- Learn basic data visualization techniques

## Not Covered

- Topics in “big data” Optimization
- Math
- Data in formats that are not tabular: ie. genetic sequences, geodatabases, images

# Why Pandas?

- Free, easy to learn
- Compatible with a growing data science ecosystem in Python
  - Visualization
  - Machine learning
  - Statistics
  - Scientific applications
- Flexible in terms of input and output



stumpy

bokkeh



pandas



# Creating DataFrame Objects

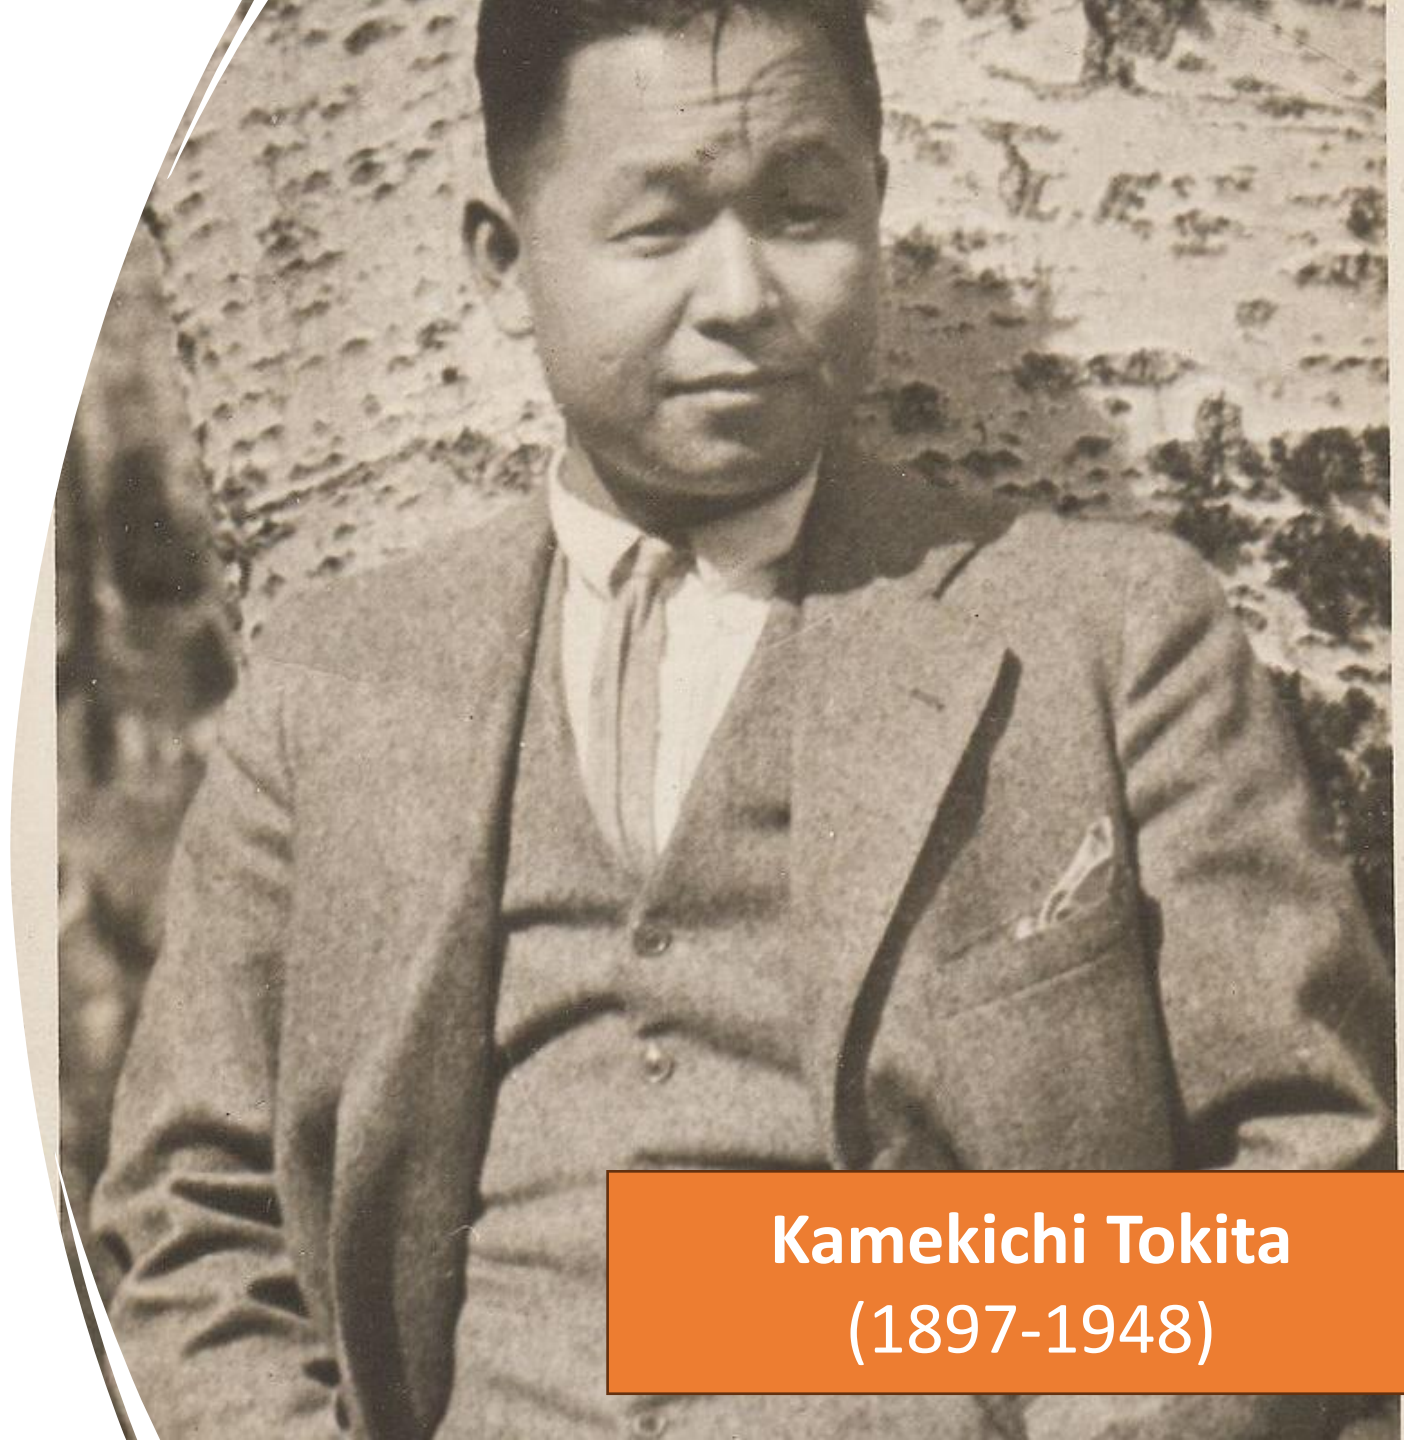
- **Can be constructed directly from hardcoded Python objects like dictionaries**
- **Loaded from a file or URL:**
  - delimited text file (.txt, .tsv)
  - comma-separated values (.csv)
  - Excel sheet (.xlsx)
  - JSON dictionary (.json)
  - Stata/SAS files (.sav, .sas, .dta)



# A Little Modern Art

---

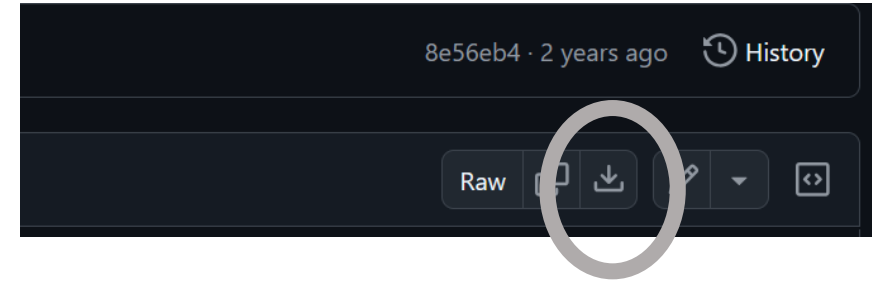
- Our first dataset is the MOMA's Watson Library Index of Asian American and Pacific Islander Artists.
- This is a collection of AAPI artists featured in the Watson Library's catalog of exhibit records and artist biographies.



**Kamekichi Tokita**  
(1897-1948)



Tokita, Kamekichi. Alley. 1929. Seattle Art Museum. Collection of Shokichi & Elise Y. Tokita.



## Loading Data From A File

- Find the download link in the *resources.md* in the class GitHub page
- Download this file from the Watson Library GitHub by clicking the download button
- Copy this .csv file into your JupyterLab project directory

# Reading Input

Reading files into Pandas is as simple as **matching the file type to the name of a function...**

- `pd.read_csv()`
- `pd.read_json()`
- `pd.read_excel()`
- `pd.read_sql()`
- `pd.read_stata()`
- ...

Format Type	Data Description	Reader	Writer
text	<a href="#">CSV</a>	<a href="#">read_csv</a>	<a href="#">to_csv</a>
text	Fixed-Width Text File	<a href="#">read_fwf</a>	NA
text	<a href="#">JSON</a>	<a href="#">read_json</a>	<a href="#">to_json</a>
text	<a href="#">HTML</a>	<a href="#">read_html</a>	<a href="#">to_html</a>
text	<a href="#">LaTeX</a>	<a href="#">Styler.to_latex</a>	NA
text	<a href="#">XML</a>	<a href="#">read_xml</a>	<a href="#">to_xml</a>
text	Local clipboard	<a href="#">read_clipboard</a>	<a href="#">to_clipboard</a>
binary	<a href="#">MS Excel</a>	<a href="#">read_excel</a>	<a href="#">to_excel</a>
binary	<a href="#">OpenDocument</a>	<a href="#">read_excel</a>	NA
binary	HDF5 Format	<a href="#">read_hdf</a>	<a href="#">to_hdf</a>

[https://pandas.pydata.org/docs/dev/user\\_guide/io.html](https://pandas.pydata.org/docs/dev/user_guide/io.html)



# Understanding Your Data (I/O)

After reading from a file to a DataFrame, check

- delimiting (are the columns separated?)
- the shape – a (row #, column #) tuple
- the types of each variable – dtypes
- which columns have nulls/nones/NaNs (and **why**)



If it is **not** what you expect, you probably need to adjust the formatting arguments passed to Pandas.

# The Pandas DataFrame

Each **DataFrame** column is a **Series**.

.index

	.columns					axis = 1		
	name	colour	location	seed	shape	sweetness	water_content	weight
0	apple	red	canada	TRUE	round	TRUE	84	100
1	banana	yellow	mexico	FALSE	long	TRUE	75	120
2	cantaloupe	orange	spain	TRUE	round	TRUE	90	1360
3	dragon fruit	magenta	china	TRUE	round	FALSE	96	600
4	elderberry	purple	austria	FALSE	round	TRUE	80	5

axis = 0

All values in the DataFrame have an indexed position in `.iloc[row_index, column_number]` form.

Total elements = `.size`  
R, C form = `.shape`

# The Pandas Series

The **Series** is the simplest of the Pandas data structures: a one-dimensional array with an index. The alpha-numeric index is exactly as long as the data.

.index	0	1	2	3	4	5	6	7
	burgundy	red	green	gray	blue	yellow	orange	teal

All values in the Series have an indexed position in  
[index] or .iloc[index\_num] form.

.index	a	b	c	d	e	f	g	h
	burgundy	red	green	gray	blue	yellow	orange	teal

# Slicing and Indexing with .loc

		Column Labels		
Index		name	colour	location
	a	apple	red	canada
	b	banana	yellow	mexico
	c	cantaloupe	orange	spain
	d	dragon fruit	magenta	china
	e	elderberry	purple	austria

**.loc slicing is inclusive on both ends because it is intended to behave like R (not Python)**

## **.loc – index and column based slicing**

`.loc['a']` returns a **pandas.Series** with the first row

`.loc['a', 'location']` returns the string “canada”

`.loc['b':'d', 'name':'colour']` returns the DataFrame below

	<b>name</b>	<b>colour</b>
b	banana	yellow
c	cantaloupe	orange
d	dragon fruit	magenta

# Slicing and Indexing with .iloc

		Column Labels		
Index		name	colour	location
	a	apple	red	canada
	b	banana	yellow	mexico
	c	cantaloupe	orange	spain
	d	dragon fruit	magenta	china
	e	elderberry	purple	austria

**.iloc slices behave like Python list slices (exclusive at the end of a range)**

## **.iloc – integer-based slicing**

`.iloc[0]` returns a **pandas.Series** with the first row

`.iloc[1, 2]` returns the string “mexico”

`.iloc[2:, 1:3]` returns the DataFrame below

	<b>colour</b>	<b>location</b>
c	orange	spain
c	magenta	china
d	purple	austria



# Boolean Indexing

- Pandas .loc and .iloc can return subsets of a DataFrame specified by a boolean array.
  - Boolean arrays filter DataFrames or Series by testing values against a condition. *NaNs evaluate to False*.

```
under21 = students['age'] < 21
```

a series of size (# of rows in students) a pd.DataFrame numeric column the “test”, must evaluate to True or False

```
studentsUnder21 = students[under21]
```

only the rows in the students DataFrame where the column 'age' < 21

`x['col_name']` is shorthand for `x.loc[:, 'col_name']`

# Boolean Indexing with Multiple Conditions

- We can combine boolean indexing with multiple conditions as follows:

```
ok_to_drink = (students['age'] >= 21) & (students['has_id'] == 'Yes')
```

a pd.DataFrame   column   test   boolean operator

a series of  
size (# of  
rows in  
students)

```
studentDrinkers = students[ok_to_drink]
```

a DataFrame with only the rows in the  
students DataFrame where the column 'age'  
> 21 AND column 'has\_id' has a 'Yes' value

# What is a NaN? (null, None, etc.)

The np.NaN is the way Pandas represents **missing values** by default, but missing values occur in almost **all** domains:

- Product of the data entry process
- Missing from the source data
- Measurement error
- Participant declined to respond
- Something went wrong computationally when creating the data