

- JupyterLab
- Jupyter Notebook
  - or an IDE of your choice...

Please type and execute this code:

```
import pandas as pd  
print(pd.__version__)
```

'2.2.3'

← you should get something like this

The Jupyter logo features the word "jupyter" in a lowercase, grey, sans-serif font. It is centered within a large orange circle that is formed by two thick, curved orange lines. Four grey circles of varying sizes are positioned around the orange circle, resembling orbits or data points.

jupyter

Raise your hand if you cannot import Pandas **or if you get a version < 2.0**

# Roll Call

- Your name
- Role or class year
- Department or major
- A source or form of tabular data you regularly use or want to learn to use

ICPSR

United States<sup>®</sup>  
**Census**  
Bureau

eurostat 



# Workshop Policies

Lecture Notes  
(.ipynb)

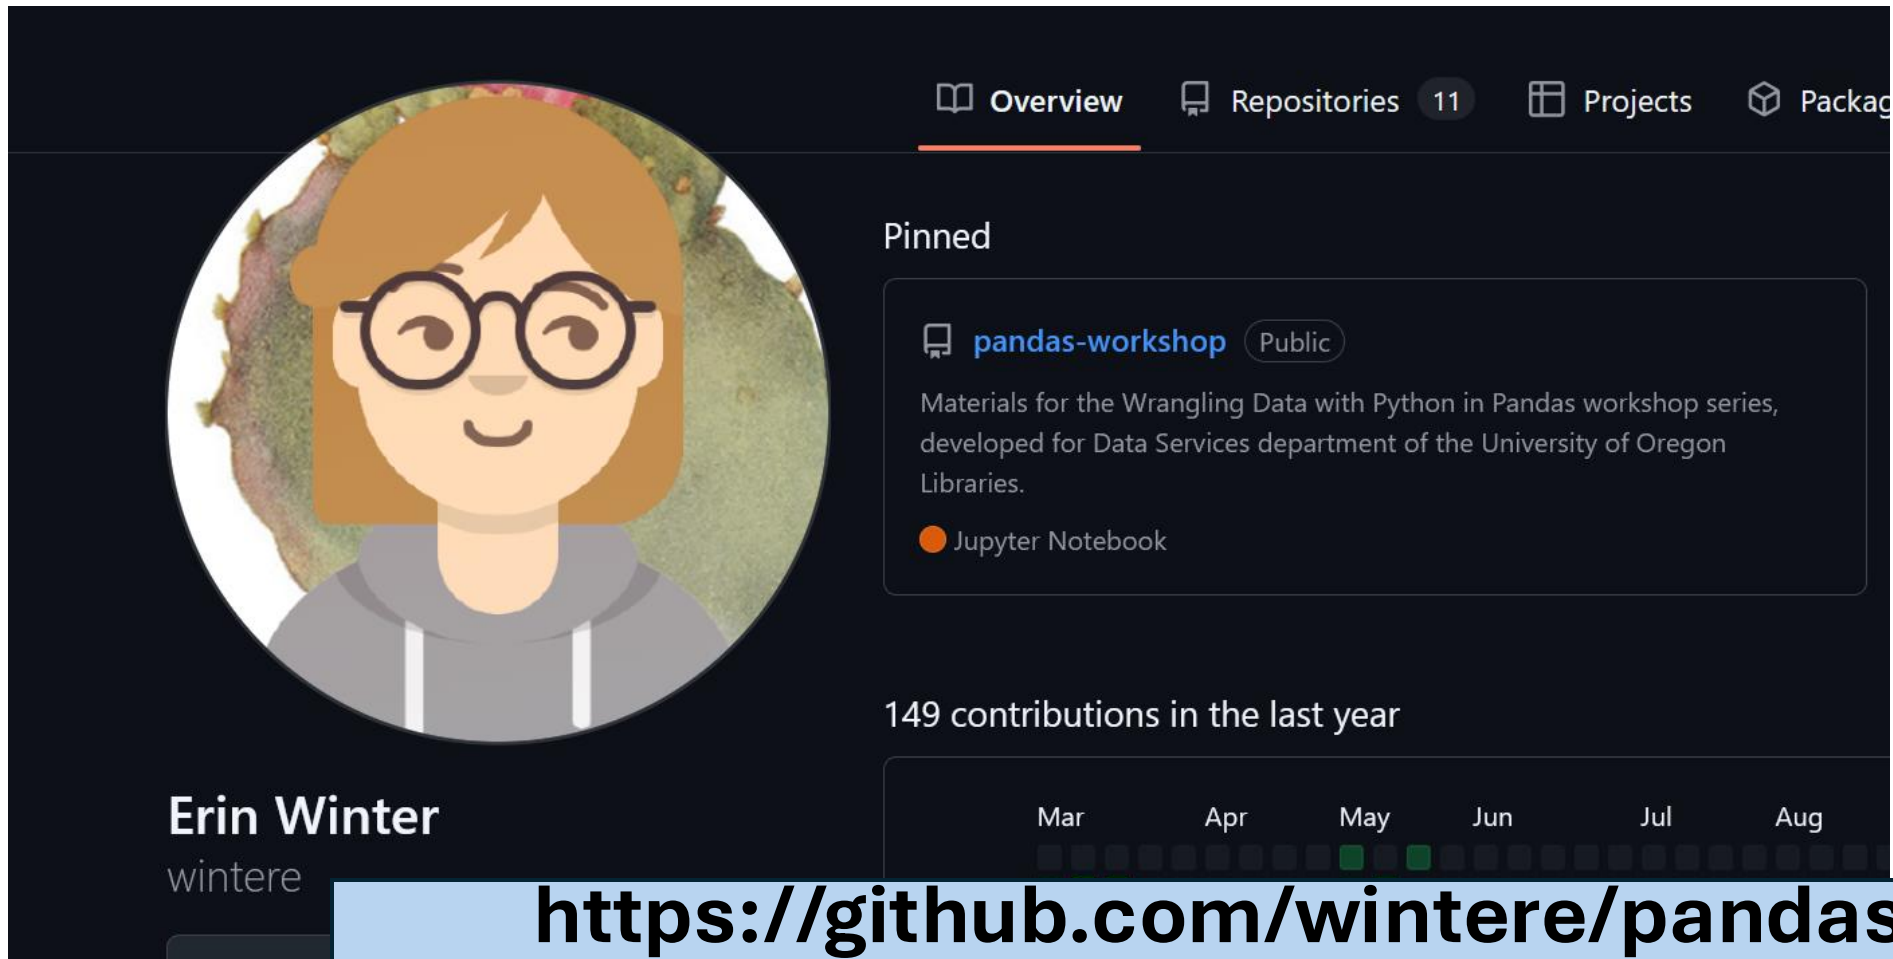
Slides

Updated a few days  
after the workshop



Link automatically  
sent to the *email you*  
*registered with* 24  
hours before the  
workshop

# Lecture Notes and Slides on GitHub



The screenshot shows the GitHub profile of Erin Winter. On the left is a circular profile picture of a woman with brown hair and glasses. Below it, the name "Erin Winter" and the username "wintere" are visible. To the right, the navigation bar includes "Overview" (selected), "Repositories" (11), "Projects", and "Packages". Under the "Pinned" section, the repository "pandas-workshop" is listed as "Public". Its description reads: "Materials for the Wrangling Data with Python in Pandas workshop series, developed for Data Services department of the University of Oregon Libraries." Below this, it is noted as a "Jupyter Notebook". Further down, it states "149 contributions in the last year" and shows a calendar view with activity bars for March through August.

<https://github.com/wintere/pandas-workshop>

# Prerequisites

- Software Carpentries:  
Plotting and Programming  
in Python

OR

- **variables, assignment**
- standard Python primitives
  - int, float, string, boolean
- control flow: if and else
- *functions, scope*



You will get opportunities to practice these if you're still grasping them.

# Workshop Objectives

## Goals

- Manipulate, analyze, and visualize data in Pandas and related libraries
- Import data from a variety of sources
- Identify and correct common problems in tabular data
- Learn basic data visualization techniques

## Not Covered

- Topics in “big data” Optimization
- Math
- Data in formats that are not tabular: ie. genetic sequences, geodatabases, images

# Why Pandas?

- Free, easy to learn
- Compatible with a growing data science ecosystem in Python
  - Visualization
  - Machine learning
  - Statistics
  - Scientific applications
- Flexible in terms of input and output




# Creating DataFrame Objects

- **Can be constructed directly from hardcoded Python objects like dictionaries**
- **Loaded from a file or URL:**
  - delimited text file (.txt, .tsv)
  - comma-separated values (.csv)
  - Excel sheet (.xlsx)
  - JSON dictionary (.json)
  - Stata/SAS files (.sav, .sas, .dta)



# A Little Modern Art

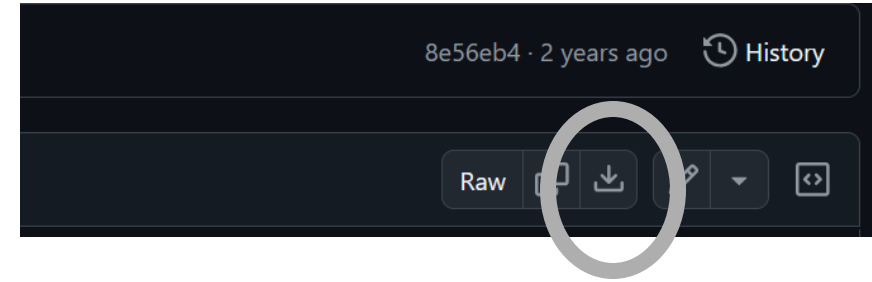
- Our first dataset is the MOMA's Watson Library Index of Asian American and Pacific Islander Artists.
- This is a collection of AAPI artists featured in the Watson Library's catalog of exhibit records and artist biographies.



**Kamekichi Tokita**  
(1897-1948)



Tokita, Kamekichi. Alley. 1929. Seattle Art Museum. Collection of Shokichi & Elise Y. Tokita.



## Loading Data From A File

- Find the download link in the *resources.md* in the class GitHub page
- Download this file from the Watson Library GitHub by clicking the download button
- Copy this .csv file into your JupyterLab project directory

# Reading Input

Reading files into Pandas is as simple as **matching the file type to the name of a function...**

- `pd.read_csv()`
- `pd.read_json()`
- `pd.read_excel()`
- `pd.read_sql()`
- `pd.read_stata()`
- ...

Format Type	Data Description	Reader	Writer
text	<a href="#">CSV</a>	<a href="#">read_csv</a>	<a href="#">to_csv</a>
text	Fixed-Width Text File	<a href="#">read_fwf</a>	NA
text	<a href="#">JSON</a>	<a href="#">read_json</a>	<a href="#">to_json</a>
text	<a href="#">HTML</a>	<a href="#">read_html</a>	<a href="#">to_html</a>
text	<a href="#">LaTeX</a>	<a href="#">Styler.to_latex</a>	NA
text	<a href="#">XML</a>	<a href="#">read_xml</a>	<a href="#">to_xml</a>
text	Local clipboard	<a href="#">read_clipboard</a>	<a href="#">to_clipboard</a>
binary	<a href="#">MS Excel</a>	<a href="#">read_excel</a>	<a href="#">to_excel</a>
binary	<a href="#">OpenDocument</a>	<a href="#">read_excel</a>	NA
binary	<a href="#">HDF5 Format</a>	<a href="#">read_hdf</a>	<a href="#">to_hdf</a>

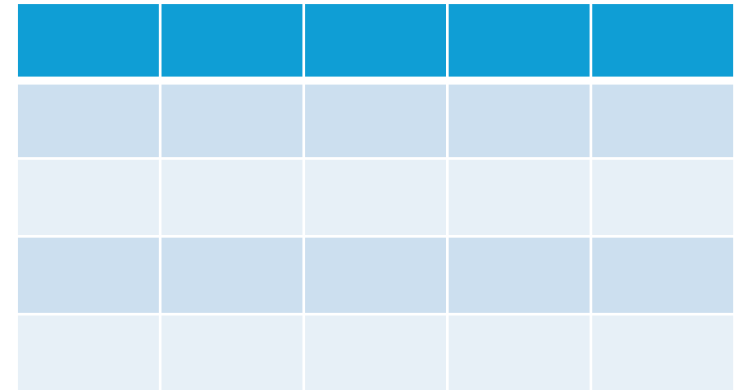
[https://pandas.pydata.org/docs/dev/user\\_guide/io.html](https://pandas.pydata.org/docs/dev/user_guide/io.html)



# Understanding Your Data (I/O)

After reading from a file to a DataFrame, check

- delimiting (are the columns separated?)
- the shape – a (row #, column #) tuple
- the types of each variable – dtypes
- which columns have nulls/nones/NaNs (and **why**)




If it is **not** what you expect, you probably need to adjust the formatting arguments passed to Pandas.

# The Pandas DataFrame

Each **DataFrame** column is a **Series**.

.index

0

1

2

3

4

.columns

axis = 1

name	colour	location	seed	shape	sweetness	water_content	weight
apple	red	canada	TRUE	round	TRUE	84	100
banana	yellow	mexico	FALSE	long	TRUE	75	120
cantaloupe	orange	spain	TRUE	round	TRUE	90	1360
dragon fruit	magenta	china	TRUE	round	FALSE	96	600
elderberry	purple	austria	FALSE	round	TRUE	80	5

axis = 0

All values in the DataFrame have an indexed position in `.iloc[row_index, column_number]` form.

Total elements = `.size`  
R, C form = `.shape`

# The Pandas Series

The **Series** is the simplest of the Pandas data structures: a one-dimensional array with an index. The alpha-numeric index is exactly as long as the data.

.index	0	1	2	3	4	5	6	7
	burgundy	red	green	gray	blue	yellow	orange	teal

All values in the Series have an indexed position in  
[index] or .iloc[index\_num] form.

.index	a	b	c	d	e	f	g	h
	burgundy	red	green	gray	blue	yellow	orange	teal

# Slicing and Indexing with .loc

		Column Labels		
Index		name	colour	location
	a	apple	red	canada
	b	banana	yellow	mexico
	c	cantaloupe	orange	spain
	d	dragon fruit	magenta	china
	e	elderberry	purple	austria

**.loc slicing is inclusive on both ends because it is intended to behave like R (not Python)**

## **.loc – index and column based slicing**

`.loc['a']` returns a **pandas.Series** with the first row

`.loc['a', 'location']` returns the string “canada”

`.loc['b':'d', 'name': 'colour']` returns the DataFrame below

	name	colour
b	banana	yellow
c	cantaloupe	orange
d	dragon fruit	magenta

# Slicing and Indexing with .iloc

		Column Labels		
Index		name	colour	location
	a	apple	red	canada
	b	banana	yellow	mexico
	c	cantaloupe	orange	spain
	d	dragon fruit	magenta	china
	e	elderberry	purple	austria

**.iloc slices behave like Python list slices (exclusive at the end of a range)**

## **.iloc – integer-based slicing**

`.iloc[0]` returns a **pandas.Series** with the first row

`.iloc[1, 2]` returns the string “mexico”

`.iloc[2:, 1:3]` returns the DataFrame below

	colour	location
c	orange	spain
c	magenta	china
d	purple	austria



# Boolean Indexing

- Pandas .loc and .iloc can return subsets of a DataFrame specified by a boolean array.
  - Boolean arrays filter DataFrames or Series by testing values against a condition. *NaNs evaluate to False*.

```
under21 = students['age'] < 21
```

a series of size (# of rows in students) a pd.DataFrame numeric column the “test”, must evaluate to True or False

```
studentsUnder21 = students[under21]
```

only the rows in the students DataFrame where the column ‘age’ < 21

`x['col_name']` is shorthand for `x.loc[:, 'col_name']`

# Boolean Indexing with Multiple Conditions

- We can combine boolean indexing with multiple conditions as follows:

```
ok_to_drink = (students['age'] >= 21) & (
    students['has_id'] == 'Yes')
```

a pd.DataFrame    column    test    boolean operator

test

a series of  
size (# of  
rows in  
students)

```
studentDrinkers = students[ok_to_drink]
```

a DataFrame with only the rows in the  
students DataFrame where the column 'age'  
> 21 AND column 'has\_id' has a 'Yes' value